

Questions

February 8, 2022

1 Data Science Challenge

```
[466]: # To install packages that are not installed by default, uncomment the last two
      ↪ lines
      # of this cell and replace <package list> with a list of necessary packages.
      # This will ensure the notebook has all the dependencies and works everywhere.

import sys
!{sys.executable} -m pip install xgboost
!{sys.executable} -m pip install plot-metric
```

Requirement already satisfied: xgboost in /opt/conda/lib/python3.7/site-packages (1.5.2)

Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.19.2)

Requirement already satisfied: scipy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.3.3)

Collecting plot-metric

Downloading plot_metric-0.0.6-py3-none-any.whl (13 kB)

Requirement already satisfied: matplotlib>=3.0.2 in /opt/conda/lib/python3.7/site-packages (from plot-metric) (3.1.3)

Requirement already satisfied: pandas>=0.23.4 in /opt/conda/lib/python3.7/site-packages (from plot-metric) (0.25.3)

Requirement already satisfied: scipy>=1.1.0 in /opt/conda/lib/python3.7/site-packages (from plot-metric) (1.3.3)

Requirement already satisfied: numpy>=1.15.4 in /opt/conda/lib/python3.7/site-packages (from plot-metric) (1.19.2)

Collecting colorlover>=0.3.0

Downloading colorlover-0.3.0-py3-none-any.whl (8.9 kB)

Requirement already satisfied: scikit-learn>=0.21.2 in /opt/conda/lib/python3.7/site-packages (from plot-metric) (1.0.2)

Requirement already satisfied: seaborn>=0.9.0 in /opt/conda/lib/python3.7/site-packages (from plot-metric) (0.9.0)

Requirement already satisfied: cycycler>=0.10 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=3.0.2->plot-metric) (0.10.0)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib>=3.0.2->plot-metric)

```

(2.4.7)
Requirement already satisfied: python-dateutil>=2.1 in
/opt/conda/lib/python3.7/site-packages (from matplotlib>=3.0.2->plot-metric)
(2.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in
/opt/conda/lib/python3.7/site-packages (from matplotlib>=3.0.2->plot-metric)
(1.3.0)
Requirement already satisfied: pytz>=2017.2 in /opt/conda/lib/python3.7/site-
packages (from pandas>=0.23.4->plot-metric) (2020.1)
Requirement already satisfied: joblib>=0.11 in /opt/conda/lib/python3.7/site-
packages (from scikit-learn>=0.21.2->plot-metric) (0.17.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.21.2->plot-metric)
(3.1.0)
Requirement already satisfied: six in /opt/conda/lib/python3.7/site-packages
(from cycloper>=0.10->matplotlib>=3.0.2->plot-metric) (1.15.0)
Installing collected packages: colorlover, plot-metric
Successfully installed colorlover-0.3.0 plot-metric-0.0.6

```

```

[469]: #Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier, plot_importance
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from plot_metric.functions import BinaryClassification
from sklearn.metrics import roc_auc_score
pd.set_option("display.max_columns", 101)

```

1.1 Data Description

Column	Description
id	The unique ID assigned to every hotel.
region	The region in which the hotel is located..
latitude	The latitude of the hotel.
longitude	The longitude of the hotel.

Column	Description
accommodation_type	The type of accommodation offered by the hotel. For example: Private room, Entire house/apt, etc.
cost	The cost of booking the hotel for one night. (in \$\$)
minimum_nights	The minimum number of nights stay required.
number_of_reviews	The number of reviews accumulated by the hotel.
reviews_per_month	The average number of reviews received by the hotel per month.
owner_id	The unique ID assigned to every owner. An owner can own multiple hotels.
owned_hotels	The number of hotels owned by the owner.
yearly_availability	It indicates if the hotel accepts bookings around the year. Values are 0 (not available for 365 days in a year) and 1 (available for 365 days in a year).

1.2 Data Wrangling & Visualization

```
[646]: # Dataset is already loaded below
data = pd.read_csv("train.csv")
print(len(data))
```

2870

There is high correlation between owner_id and id, and number_of_reviews and reviews_per_month. One of these columns will be removed as explained below.

```
[647]: cor_plot = data.corr()
cor_plot.style.background_gradient(cmap = 'coolwarm')
```

```
[647]: <pandas.io.formats.style.Styler at 0x7f3d4021a790>
```

2 Histogram plot of numerical columns


```
[648]: plt.hist(data['latitude'], density=True, bins = 30)
```

```
[648]: (array([0.02668955, 0.          , 0.          , 0.02668955, 0.08006865,
0.34696413, 0.40034323, 0.13344774, 0.53379097, 1.12096104,
1.94833704, 2.56219666, 3.44295176, 6.72576622, 6.69907668,
7.33962584, 7.28624674, 5.55142609, 8.19369139, 8.72748236,
3.12267718, 1.57468336, 3.38957266, 2.24192207, 1.8682684 ,
```

```

1.01420284, 0.90744465, 0.90744465, 0.32027458, 0.10675819]],
array([40.50708 , 40.520135, 40.53319 , 40.546245, 40.5593 , 40.572355,
40.58541 , 40.598465, 40.61152 , 40.624575, 40.63763 , 40.650685,
40.66374 , 40.676795, 40.68985 , 40.702905, 40.71596 , 40.729015,
40.74207 , 40.755125, 40.76818 , 40.781235, 40.79429 , 40.807345,
40.8204 , 40.833455, 40.84651 , 40.859565, 40.87262 , 40.885675,
40.89873 ]),
<a list of 30 Patch objects>)

```



```


[649]: plt.hist(data['longitude'] , density=True, bins = 30)

```

```


[649]: (array([ 0.02005865,  0.          ,  0.          ,  0.02005865,  0.06017594,
 0.          ,  0.04011729,  0.04011729,  0.          ,  0.08023458,
 0.02005865,  0.02005865,  0.2607624 ,  3.83120141, 11.19272454,
 7.7827547 , 11.99507039,  8.92609753,  5.11495476,  2.74803452,
 0.88258043,  1.14334283,  0.5215248 ,  0.60175938,  0.76222855,
 0.58170074,  0.22064511,  0.2607624 ,  0.28082105,  0.16046917]),
array([-74.24285 , -74.22547933, -74.20810867, -74.190738 ,
-74.17336733, -74.15599667, -74.138626 , -74.12125533,
-74.10388467, -74.086514 , -74.06914333, -74.05177267,
-74.034402 , -74.01703133, -73.99966067, -73.98229 ,
-73.96491933, -73.94754867, -73.930178 , -73.91280733,
-73.89543667, -73.878066 , -73.86069533, -73.84332467,
-73.825954 , -73.80858333, -73.79121267, -73.773842 ,
-73.75647133, -73.73910067, -73.72173 ]),
<a list of 30 Patch objects>)

```




```
[650]: plt.hist(data['cost'] , density=True, bins = 30)
```

```
[650]: (array([2.72913445e-03, 1.63245772e-04, 6.80190719e-05, 1.04644726e-05,
 7.32513081e-06, 9.41802533e-06, 1.04644726e-06, 2.09289452e-06,
 3.13934178e-06, 1.04644726e-06, 0.00000000e+00, 0.00000000e+00,
 1.04644726e-06, 2.09289452e-06, 1.04644726e-06, 0.00000000e+00,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 1.04644726e-06,
 1.04644726e-06, 0.00000000e+00, 0.00000000e+00, 1.04644726e-06,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
 0.00000000e+00, 1.04644726e-06]),
array([ 10.          , 342.96666667, 675.93333333, 1008.9          ,
 1341.86666667, 1674.83333333, 2007.8          , 2340.76666667,
 2673.73333333, 3006.7          , 3339.66666667, 3672.63333333,
 4005.6          , 4338.56666667, 4671.53333333, 5004.5          ,
 5337.46666667, 5670.43333333, 6003.4          , 6336.36666667,
 6669.33333333, 7002.3          , 7335.26666667, 7668.23333333,
 8001.2          , 8334.16666667, 8667.13333333, 9000.1          ,
 9333.06666667, 9666.03333333, 9999.          ]),
<a list of 30 Patch objects>)
```




```
[651]: plt.hist(data['minimum_nights'] , density=True, bins = 30)
```

```
[651]: (array([2.92850509e-02, 1.99004280e-04, 1.88530371e-04, 2.09478190e-05,
 1.04739095e-05, 1.57108642e-04, 2.09478190e-05, 0.00000000e+00,
 2.09478190e-05, 0.00000000e+00, 1.15213004e-04, 1.04739095e-05,
 0.00000000e+00, 0.00000000e+00, 1.04739095e-05, 1.04739095e-05,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
 0.00000000e+00, 1.04739095e-05]),
array([ 1.          , 34.26666667, 67.53333333, 100.8          ,
 134.06666667, 167.33333333, 200.6          , 233.86666667,
 267.13333333, 300.4          , 333.66666667, 366.93333333,
 400.2          , 433.46666667, 466.73333333, 500.          ,
 533.26666667, 566.53333333, 599.8          , 633.06666667,
 666.33333333, 699.6          , 732.86666667, 766.13333333,
 799.4          , 832.66666667, 865.93333333, 899.2          ,
 932.46666667, 965.73333333, 999.          ]),
<a list of 30 Patch objects>)
```




```
[652]: plt.hist(data['number_of_reviews'] , density=True, bins = 30)
```

```
[652]: (array([5.39849160e-02, 9.05041238e-03, 4.26057425e-03, 2.61985622e-03,
 1.56132845e-03, 1.16438054e-03, 5.82190270e-04, 5.29263882e-04,
 3.96947912e-04, 2.64631941e-04, 3.44021523e-04, 2.64631941e-04,
 1.58779165e-04, 2.11705553e-04, 7.93895823e-05, 1.85242359e-04,
 5.29263882e-05, 5.29263882e-05, 7.93895823e-05, 2.64631941e-05,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 2.64631941e-05,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 2.64631941e-05,
 0.00000000e+00, 2.64631941e-05]),
array([ 0.          , 13.16666667, 26.33333333, 39.5          ,
 52.66666667, 65.83333333, 79.          , 92.16666667,
105.33333333, 118.5          , 131.66666667, 144.83333333,
158.          , 171.16666667, 184.33333333, 197.5          ,
210.66666667, 223.83333333, 237.          , 250.16666667,
263.33333333, 276.5          , 289.66666667, 302.83333333,
316.          , 329.16666667, 342.33333333, 355.5          ,
368.66666667, 381.83333333, 395.          ]),
<a list of 30 Patch objects>)
```



```
[653]: plt.hist(data['owned_hotels'] , density=True, bins = 30)
```


```
[653]: (array([8.21487356e-02, 2.30862957e-03, 1.44289348e-03, 6.41285992e-04,
 2.11624377e-03, 9.29864689e-04, 0.00000000e+00, 6.41285992e-04,
 2.56514397e-04, 3.20642996e-05, 0.00000000e+00, 9.93993288e-04,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
 0.00000000e+00, 2.88578696e-04, 0.00000000e+00, 0.00000000e+00,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
 0.00000000e+00, 2.24450097e-04]),
array([ 1.          , 11.86666667, 22.73333333, 33.6          ,
 44.46666667, 55.33333333, 66.2          , 77.06666667,
 87.93333333, 98.8          , 109.66666667, 120.53333333,
131.4          , 142.26666667, 153.13333333, 164.          ,
174.86666667, 185.73333333, 196.6          , 207.46666667,
218.33333333, 229.2          , 240.06666667, 250.93333333,
261.8          , 272.66666667, 283.53333333, 294.4          ,
305.26666667, 316.13333333, 327.          ]),
<a list of 30 Patch objects>)
```

Region wise count

```
[654]: data['region'].value_counts().plot(kind = 'barh')
```


```
[654]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3d3a17cc50>
```



Accommodation Type wise count

```
[655]: data['accommodation_type'].value_counts().plot(kind = 'barh')
```


```
[655]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3d3a0a6f10>
```



Cost across regions

```
[656]: data.groupby(['region'])['cost'].mean().plot(kind = 'barh')
```


```
[656]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3d3a027710>
```



Cost across Accomodation Type

```
[657]: data.groupby(['accommodation_type'])['cost'].mean().plot(kind = 'barh')
```

```
[657]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3d39f941d0>
```



Since percentage of unique values in 'id', and 'owner_id' column is more than 80%, I am dropping

the columns, since they won't add much value to the model

```
[658]: print(data['id'].nunique()/len(data))  
print(data['owner_id'].nunique()/len(data))
```

```
1.0  
0.8261324041811847
```

```
[659]: del data['id']  
del data['owner_id']
```

```
[660]: set(data['accommodation_type'])
```


```
[660]: {'Entire home/apt', 'Private room', 'Shared room'}
```

Plot to check number of classes. The data is well sampled in this case

```
[661]: print(data['yearly_availability'].value_counts())  
data['yearly_availability'].value_counts().plot(kind = 'barh')
```

```
0    1439  
1    1431  
Name: yearly_availability, dtype: int64
```

```
[661]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3d39f5ea90>
```



```
[662]: data['reviews_per_month'].value_counts()
```

```
[662]: 1.00    76
        0.16    40
        0.11    34
        0.05    32
        0.12    30
        ..
        2.70     1
        3.04     1
        6.59     1
        5.92     1
        5.97     1
        Name: reviews_per_month, Length: 419, dtype: int64
```

Checking if data contains null values. Dropping the column reviews_per_month, since it has more than 20% missing values and it is correlated with number_of_reviews as shown earlier.

```
[663]: data.isnull().sum()/len(data)
```

```
[663]: region                0.00000
        latitude             0.00000
        longitude            0.00000
        accommodation_type    0.00000
        cost                  0.00000
        minimum_nights        0.00000
        number_of_reviews      0.00000
        reviews_per_month     0.23554
        owned_hotels           0.00000
        yearly_availability     0.00000
        dtype: float64
```

```
[541]: del data['reviews_per_month']
```

Getting the numerical columns

```
[542]: num_cols = list(data.columns[data.dtypes.apply(lambda c: np.issubdtype(c, np.
↪number))])
        num_cols.remove('yearly_availability')
        num_cols
```

```
[542]: ['latitude',
        'longitude',
        'cost',
        'minimum_nights',
        'number_of_reviews',
        'owned_hotels']
```

Label Encoder vs One hot Encoder : Label Encoder might generate a bias in the data, which is unwarranted. And since one hot encoder did not blow up the data size, going forward with One

hot Encoder

```
[543]: # le = LabelEncoder()  
# data['region'] = le.fit_transform(data['region'])  
# data['accommodation_type'] = le.fit_transform(data['accommodation_type'])
```

```
[544]: data.columns
```

```
[544]: Index(['region', 'latitude', 'longitude', 'accommodation_type', 'cost',  
        'minimum_nights', 'number_of_reviews', 'owned_hotels',  
        'yearly_availability'],  
        dtype='object')
```

```
[545]: data = pd.get_dummies(data, columns =  
    ↳ ['region', 'accommodation_type'], drop_first = True)
```

```
[546]: data.columns
```

```
[546]: Index(['latitude', 'longitude', 'cost', 'minimum_nights', 'number_of_reviews',  
        'owned_hotels', 'yearly_availability', 'region_Brooklyn',  
        'region_Manhattan', 'region_Queens', 'region_Staten Island',  
        'accommodation_type_Private room', 'accommodation_type_Shared room'],  
        dtype='object')
```

```
[547]: data.shape
```

```
[547]: (2870, 13)
```

```
[548]: data.head()
```

```
[548]:
```

	latitude	longitude	cost	minimum_nights	number_of_reviews	owned_hotels	\
0	40.71854	-74.00439	170	5	7	1	
1	40.64446	-73.95030	65	3	238	1	
2	40.78573	-73.81062	85	1	0	1	
3	40.73863	-73.98002	210	30	0	65	
4	40.82426	-73.94630	75	3	38	3	

	yearly_availability	region_Brooklyn	region_Manhattan	region_Queens	\
0	0	0	1	0	
1	0	1	0	0	
2	1	0	0	1	
3	1	0	1	0	
4	1	0	1	0	

	region_Staten Island	accommodation_type_Private room	\
0	0	0	
1	0	0	
2	0	1	

3	0	1
4	0	0

accommodation_type_Shared room	
0	0
1	0
2	0
3	0
4	1

```
[549]: #Explore columns
data.columns
```

```
[549]: Index(['latitude', 'longitude', 'cost', 'minimum_nights', 'number_of_reviews',
         'owned_hotels', 'yearly_availability', 'region_Brooklyn',
         'region_Manhattan', 'region_Queens', 'region_Staten Island',
         'accommodation_type_Private room', 'accommodation_type_Shared room'],
        dtype='object')
```

```
[550]: #Description
data.describe()
```

```
[550]:
```

	latitude	longitude	cost	minimum_nights	\
count	2870.000000	2870.000000	2870.000000	2870.000000	
mean	40.731224	-73.950158	195.943206	11.530314	
std	0.054942	0.049745	406.184714	37.972339	
min	40.507080	-74.242850	10.000000	1.000000	
25%	40.692462	-73.984003	75.000000	1.000000	
50%	40.728250	-73.956720	120.000000	3.000000	
75%	40.762658	-73.934202	200.000000	6.000000	
max	40.898730	-73.721730	9999.000000	999.000000	

	number_of_reviews	owned_hotels	yearly_availability	region_Brooklyn	\
count	2870.000000	2870.000000	2870.000000	2870.000000	
mean	16.315331	8.411498	0.498606	0.374564	
std	32.481722	27.105522	0.500085	0.484095	
min	0.000000	1.000000	0.000000	0.000000	
25%	1.000000	1.000000	0.000000	0.000000	
50%	4.000000	1.000000	0.000000	0.000000	
75%	16.000000	3.000000	1.000000	1.000000	
max	395.000000	327.000000	1.000000	1.000000	

	region_Manhattan	region_Queens	region_Staten Island	\
count	2870.000000	2870.000000	2870.000000	
mean	0.464460	0.12892	0.004878	
std	0.498822	0.33517	0.069685	
min	0.000000	0.000000	0.000000	

25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000

	accommodation_type_Private room	accommodation_type_Shared room
count	2870.000000	2870.000000
mean	0.426829	0.224042
std	0.494703	0.417022
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	1.000000	0.000000
max	1.000000	1.000000

Droppping duplicate entries, if any

```
[551]: data = data.drop_duplicates()
data.shape
```

```
[551]: (2870, 13)
```

```
[552]: data.columns
```

```
[552]: Index(['latitude', 'longitude', 'cost', 'minimum_nights', 'number_of_reviews',
          'owned_hotels', 'yearly_availability', 'region_Brooklyn',
          'region_Manhattan', 'region_Queens', 'region_Staten Island',
          'accommodation_type_Private room', 'accommodation_type_Shared room'],
          dtype='object')
```

```
[554]: yearly_availability = data['yearly_availability']
del data['yearly_availability']
data['yearly_availability'] = yearly_availability
```

```
[555]: data
```

```
[555]:
```

	latitude	longitude	cost	minimum_nights	number_of_reviews	\
0	40.71854	-74.00439	170	5	7	
1	40.64446	-73.95030	65	3	238	
2	40.78573	-73.81062	85	1	0	
3	40.73863	-73.98002	210	30	0	
4	40.82426	-73.94630	75	3	38	
...	
2865	40.74316	-73.98038	400	2	0	
2866	40.73523	-73.99465	180	3	2	
2867	40.76619	-73.98987	179	3	17	
2868	40.74637	-73.97207	200	30	0	


```
2869 40.79208 -73.96482 1000          30          24
```

```

      owned_hotels  region_Brooklyn  region_Manhattan  region_Queens  \
0                1                0                1                0
1                1                1                0                0
2                1                0                0                1
3               65                0                1                0
4                3                0                1                0
...
2865              1                0                1                0
2866              1                0                1                0
2867              1                0                1                0
2868             49                0                1                0
2869             11                0                1                0

```

```

      region_Staten Island  accommodation_type_Private room  \
0                        0                                0
1                        0                                0
2                        0                                1
3                        0                                1
4                        0                                0
...
2865                    0                                1
2866                    0                                1
2867                    0                                0
2868                    0                                1
2869                    0                                0

```

```

      accommodation_type_Shared room  yearly_availability
0                                0                        0
1                                0                        0
2                                0                        1
3                                0                        1
4                                1                        1
...
2865                            0                        1
2866                            0                        1
2867                            0                        0
2868                            0                        1
2869                            1                        1

```

```
[2870 rows x 13 columns]
```

```
[556]: X = data.iloc[:, :-1]
      y = data.iloc[:, -1]
```

```
[557]: num_cols
```

```
[557]: ['latitude',
        'longitude',
        'cost',
        'minimum_nights',
        'number_of_reviews',
        'owned_hotels']
```

Scaling numerical columns

```
[558]: sc = StandardScaler()
X[num_cols] = sc.fit_transform(X[num_cols])
```

```
[559]: X.columns
```

```
[559]: Index(['latitude', 'longitude', 'cost', 'minimum_nights', 'number_of_reviews',
            'owned_hotels', 'region_Brooklyn', 'region_Manhattan', 'region_Queens',
            'region_Staten Island', 'accommodation_type_Private room',
            'accommodation_type_Shared room'],
            dtype='object')
```

```
[560]: X
```

```
[560]:
```

	latitude	longitude	cost	minimum_nights	number_of_reviews	\
0	-0.230897	-1.090408	-0.063882	-0.172006	-0.286837	
1	-1.579468	-0.002863	-0.322430	-0.224685	6.826094	
2	0.992246	2.805573	-0.273182	-0.277364	-0.502380	
3	0.134826	-0.600420	0.034613	0.486483	-0.502380	
4	1.693655	0.077562	-0.297806	-0.224685	0.667712	
...	
2865	0.217291	-0.607658	0.502462	-0.251024	-0.502380	
2866	0.072932	-0.894574	-0.039258	-0.224685	-0.440796	
2867	0.636535	-0.798466	-0.041720	-0.224685	0.021082	
2868	0.275727	-0.440575	0.009989	0.486483	-0.502380	
2869	1.107843	-0.294805	1.979880	0.486483	0.236626	

	owned_hotels	region_Brooklyn	region_Manhattan	region_Queens	\
0	-0.273479	0	1	0	
1	-0.273479	1	0	0	
2	-0.273479	0	0	1	
3	2.088075	0	1	0	
4	-0.199680	0	1	0	
...	
2865	-0.273479	0	1	0	
2866	-0.273479	0	1	0	
2867	-0.273479	0	1	0	
2868	1.497687	0	1	0	
2869	0.095514	0	1	0	

	region_Staten Island	accommodation_type_Private room \
0	0	0
1	0	0
2	0	1
3	0	1
4	0	0
...
2865	0	1
2866	0	1
2867	0	0
2868	0	1
2869	0	0

	accommodation_type_Shared room
0	0
1	0
2	0
3	0
4	1
...	...
2865	0
2866	0
2867	0
2868	0
2869	1

[2870 rows x 12 columns]

[561]: y

```
[561]: 0      0
      1      0
      2      1
      3      1
      4      1
      ..
      2865    1
      2866    1
      2867    0
      2868    1
      2869    1
      Name: yearly_availability, Length: 2870, dtype: int64
```

Splitting the train data into train and validation data set.

```
[562]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=123)
```

Tried Naive Bayes, Decision Tree, Random Forest and XGB classification models.

```
[563]: gnb = GaussianNB()
gnb.fit(X_train,y_train)
```

```
[563]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
[699]: y_pred = gnb.predict(X_train)
accuracy = np.sum((y_train==y_pred))/len(y_pred)
print('Accuracy_Train: Naive Bayes', accuracy*100)
```

Accuracy_Train: Naive Bayes 80.96689895470384

```
[697]: y_pred = gnb.predict(X_test)
accuracy = np.sum((y_test==y_pred))/len(y_pred)
print('Accuracy_Test: Naive Bayes', accuracy*100)
```

Accuracy_Test: Naive Bayes 79.44250871080139

```
[698]: print("The classification report is as follows...\n")
print(classification_report(y_pred,y_test))
```

The classification report is as follows...

	precision	recall	f1-score	support
0	0.91	0.74	0.81	348
1	0.68	0.88	0.77	226
accuracy			0.79	574
macro avg	0.80	0.81	0.79	574
weighted avg	0.82	0.79	0.80	574

```
[628]: cm = (confusion_matrix(y_test,y_pred))
df_cm = pd.DataFrame(cm,index=['Class 0','Class 1'], columns=['Class 0','Class_
↳ 1'])
print("Confusion matrix\n")
df_cm
```

Confusion matrix

```
[628]:      Class 0  Class 1
Class 0      261      21
```

Since train and test accuracy is the same, the model is not overfitting

```
[568]: dt = DecisionTreeClassifier(criterion='gini',random_state=123)
dt.fit(X_train,y_train)
```

```
[568]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                             max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort=False,
                             random_state=123, splitter='best')
```

```
[695]: y_pred = dt.predict(X_train)
accuracy = np.sum((y_train==y_pred))/len(y_pred)
print('Accuracy_Train: Decision Tree', accuracy*100)
```

Accuracy_Train: Decision Tree 100.0

```
[696]: y_pred = dt.predict(X_test)
accuracy = np.sum((y_test==y_pred))/len(y_pred)
print('Accuracy_Test: Decision Tree', accuracy*100)
```

Accuracy_Test: Decision Tree 88.32752613240417

```
[ ]:
```

```
[571]: rf = RandomForestClassifier(n_estimators=500,random_state =123, max_depth =15)
rf.fit(X_train, y_train)
```

```
[571]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                             max_depth=15, max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=500,
                             n_jobs=None, oob_score=False, random_state=123,
                             verbose=0, warm_start=False)
```

```
[694]: y_pred = rf.predict(X_train)
accuracy = np.sum((y_train==y_pred))/len(y_pred)
print('Accuracy_Train - Random Forest', accuracy*100)
```

Accuracy_Train - Random Forest 100.0

```
[693]: y_pred = rf.predict(X_test)
accuracy = np.sum((y_test==y_pred))/len(y_pred)
print('Accuracy_Test - Random Forest',accuracy*100)
```

Accuracy_Test - Random Forest 93.90243902439023

[]:

```
[574]: xgb = XGBClassifier(seed = 123)
       xgb.fit(X_train, y_train)
```

[19:05:42] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1]. warnings.warn(label_encoder_deprecation_msg, UserWarning)


```
[574]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                    colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
                    gamma=0, gpu_id=-1, importance_type=None,
                    interaction_constraints='', learning_rate=0.300000012,
                    max_delta_step=0, max_depth=6, min_child_weight=1, missing=nan,
                    monotone_constraints='()', n_estimators=100, n_jobs=8,
                    num_parallel_tree=1, objective='binary:logistic',
                    predictor='auto', random_state=123, reg_alpha=0, reg_lambda=1,
                    scale_pos_weight=1, seed=123, subsample=1, tree_method='exact',
                    use_label_encoder=True, validate_parameters=1, ...)
```

```
[700]: y_pred = xgb.predict(X_test)
       accuracy = np.sum((y_test==y_pred))/len(y_pred)
       print('Accuracy_Test - XGB:',accuracy*100)
```

Accuracy_Test - XGB: 93.03135888501743

```
[627]: bc = BinaryClassification(y_test, y_pred, labels=["Class 0", "Class 1"])

       # Figures
       plt.figure(figsize=(5,5))
       bc.plot_roc_curve()
       plt.show()
```



```
[577]: roc_auc_score(y_test, y_pred)
```

```
[577]: 0.93023171108714659
```

The Random Forest model performs best with almost 94% accuracy. So will use the same model for predicting values for the test data

2.1 Visualization, Modeling, Machine Learning

Build a model that categorizes hotels on the basis of their yearly availability. Identify how different features influence the decision. Please explain the findings effectively to technical and non-technical audiences using comments and visualizations, if appropriate. - **Build an optimized model that effectively solves the business problem.** - **The model will be evaluated on the basis of Accuracy.** - Read the test.csv file and prepare features for testing.

```
[664]: #Loading Test data
test=pd.read_csv('test.csv')
test.head()
```

```
[664]:
```

	id	region	latitude	longitude	accommodation_type	cost	\
0	19215	Brooklyn	40.70912	-73.94513	Shared room	135	
1	36301	Brooklyn	40.57646	-73.96641	Entire home/apt	69	
2	40566	Manhattan	40.76616	-73.98228	Private room	225	
3	33694	Manhattan	40.77668	-73.94587	Shared room	125	
4	28873	Manhattan	40.80279	-73.94450	Entire home/apt	43	

	minimum_nights	number_of_reviews	reviews_per_month	owner_id	\
0	2	22	0.66	4360212	
1	2	8	0.90	181356989	
2	30	0	NaN	13773574	
3	30	9	0.82	6788748	
4	1	13	0.72	105061915	

	owned_hotels
0	1
1	2
2	12
3	1
4	2

Doing the same preprocessing steps done for the train data

```
[665]: id = test['id']
```

```
[666]: del test['id']
del test['owner_id']
del test['reviews_per_month']
```

```
[667]: test = pd.get_dummies(test, columns =
↳ ['region', 'accommodation_type'], drop_first = True)
```

```
[668]: test.shape
```

```
[668]: (718, 12)
```

```
[669]: test = test.drop_duplicates()
test.shape
```

```
[669]: (718, 12)
```

```
[670]: test[num_cols] = sc.transform(test[num_cols])
```

```
[671]: test
```

```
[671]:
```

	latitude	longitude	cost	minimum_nights	number_of_reviews	\
0	-0.402381	0.101086	-0.150064	-0.251024	0.175042	

1	-2.817356	-0.326774	-0.312580	-0.251024	-0.256045
2	0.635989	-0.645860	0.071548	0.486483	-0.502380
3	0.827497	0.086207	-0.174688	0.486483	-0.225253
4	1.302810	0.113753	-0.376602	-0.277364	-0.102085
..
713	2.087595	0.695225	-0.292881	-0.251024	0.606129
714	-0.381628	0.054641	-0.125441	-0.224685	0.821672
715	-0.315365	-0.825207	-0.260871	-0.251024	1.345134
716	0.232401	-0.439972	0.009989	0.486483	-0.502380
717	-0.432054	-1.286243	-0.066344	0.486483	-0.440796

	owned_hotels	region_Brooklyn	region_Manhattan	region_Queens	\
0	-0.273479	1	0	0	
1	-0.236580	1	0	0	
2	0.132413	0	1	0	
3	-0.273479	0	1	0	
4	-0.236580	0	1	0	
..	
713	-0.273479	0	0	0	
714	-0.273479	1	0	0	
715	-0.273479	0	1	0	
716	0.538305	0	1	0	
717	4.154435	0	1	0	

	region_Staten Island	accommodation_type_Private room	\
0	0	0	
1	0	0	
2	0	1	
3	0	0	
4	0	0	
..	
713	0	0	
714	0	0	
715	0	0	
716	0	1	
717	0	1	

	accommodation_type_Shared room
0	1
1	0
2	0
3	1
4	0
..	...
713	0
714	0
715	0

```
716                                0
717                                0
```

```
[718 rows x 12 columns]
```

```
[672]: test_output = rf.predict(test)
```

Highlight the most important features of the model for management.

Task:


- Visualize the top 20 features and their feature importance.

```
[676]: xgb.feature_importances_
```

```
[676]: array([0.0114952 , 0.01148161, 0.01226087, 0.01423117, 0.01824551,
          0.10496806, 0.01417442, 0.00985813, 0.01053431, 0.
          0.6363641 , 0.15638658], dtype=float32)
```

```
[677]: plot_importance(xgb)
```

```
[677]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3d39f58490>
```



Task:

- **Submit the predictions on the test dataset using your optimized model** For each record in the test set (`test.csv`), predict the value of the `yearly_availability` variable. Submit a CSV file with a header row and one row per test entry.

The file (`submissions.csv`) should have exactly 2 columns: - `id` - `yearly__availability`

```
[678]: submission_df = pd.DataFrame(columns=['id', 'yearly_availability'])
```

```
[679]: test_output.shape
```

```
[679]: (718,)
```

```
[682]: submission_df['id'] = id
       submission_df['yearly_availability'] = test_output
```

```
[683]: submission_df
```

```
[683]:
```

	id	yearly_availability
0	19215	0
1	36301	0
2	40566	1
3	33694	0
4	28873	0
..
713	26801	0
714	20110	0
715	31383	0
716	47135	1
717	13154	1

```
[718 rows x 2 columns]
```

```
[684]: #Submission
       submission_df.to_csv('submissions.csv', index=False)
```